


Fundamentals Of Real-time Systems

Kizito NKURIKIYEYEU, Ph.D.

Readings

- Read chapter 1 of the textbook¹
- Skip section 1.3 on page 16-23



¹ This lecture is based on Ovaska, L. P. A. J. (2012). Real-Time Systems Design and Analysis: Tools for the Practitioner by Phillip A. Laplante (4th Edition). Wiley-IEEE Press

System

A system is a mapping of a set of inputs into a set of outputs.




FIG 1. A system with n inputs and m outputs.

- 1** It is an assembly of components connected together in an organized way
- 2** A system is fundamentally altered if a component joins or leaves it
- 3** It has a purpose
- 4** It has a degree of permanence
- 5** It has been defined as being of particular interest

Example —A real-time control system

- Inputs are excitations and outputs are corresponding responses
- Inputs and outputs may be digital or analog
- Inputs are associated with sensors, cameras, etc.
- Outputs with actuators, displays, etc.

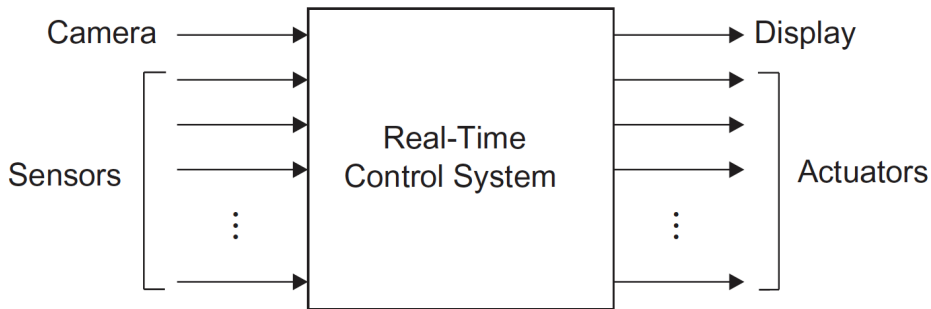


FIG 2. A real-time control system including inputs from a camera and multiple sensors, as well as outputs to a display and multiple actuators

Representation of real-time system

- A sequence of jobs to be scheduled and performance to be predicted
- Ignores the usual fact that the input sources and hardware under control may be highly complex




FIG 3. A classic representation of a real-time system as a sequence of schedulable jobs.

Response Time

The time between the presentation of a set of inputs to a system and the realization of the required behavior, including the availability of all associated outputs, is called the response time of the system

- How fast and punctual does it need to be? —Depends on the specific real-time system
- But what is a real-time system if every system somehow has deadlines and a response time?

Real-Time System

- A real-time system is a computer system that must satisfy bounded response-time constraints or risk severe consequences, including failure
- Real-time systems are defined as those systems in which the correctness of the system depends not only on the logical result of computation, but also on the time at which the results are produced.
- Real-time systems respond to external events in a timely fashion and the response time is guaranteed

Failed System

- A failed system is a system that cannot satisfy one or more of the requirements stipulated in the system requirements specification
- Hence, rigorous specification of the system operating criteria, including timing constraints, is necessary

Real-Time System

- A real-time system is a computer system that must satisfy bounded response-time constraints or risk severe consequences, including failure
- Real-time systems are defined as those systems in which the correctness of the system depends not only on the logical result of computation, but also on the time at which the results are produced.
- Real-time systems respond to external events in a timely fashion and the response time is guaranteed

Failed System

- A failed system is a system that cannot satisfy one or more of the requirements stipulated in the system requirements specification
- Hence, rigorous specification of the system operating criteria, including timing constraints, is necessary

Real-time system

- Responding to external events includes
 - Recognize when an event occurs
 - Perform the required processing
 - Output the result within a given time constraint
- Timing constraints include
 - Finish time
 - Both start time and finish time
- Example: DVD player
 - Controlling system —DVD player must decode both the video and audio streams from the disc simultaneously
 - Controlled system —Remote control is viewed as a sensor to feed pause and language selection events into DVD player

Real-time system

- Responding to external events includes
 - Recognize when an event occurs
 - Perform the required processing
 - Output the result within a given time constraint
- Timing constraints include
 - Finish time
 - Both start time and finish time
- Example: DVD player
 - Controlling system —DVD player must decode both the video and audio streams from the disc simultaneously
 - Controlled system —Remote control is viewed as a sensor to feed pause and language selection events into DVD player

Real-time system

- Responding to external events includes
 - Recognize when an event occurs
 - Perform the required processing
 - Output the result within a given time constraint
- Timing constraints include
 - Finish time
 - Both start time and finish time
- Example: DVD player
 - Controlling system —DVD player must decode both the video and audio streams from the disc simultaneously
 - Controlled system —Remote control is viewed as a sensor to feed pause and language selection events into DVD player

Degrees of “Real-Time”

Soft Real-Time System

A soft real-time system is one in which performance is degraded but not destroyed by failure to meet response-time constraints

Hard Real-Time System

A hard real-time system is one in which failure to meet even a single deadline may lead to complete or catastrophic system failure.

Soft Real-Time System

A firm real-time system is one in which a few missed deadlines will not lead to total failure, but missing more than a few may lead to complete or catastrophic system failure

Degrees of “Real-Time”

Soft Real-Time System

A soft real-time system is one in which performance is degraded but not destroyed by failure to meet response-time constraints

Hard Real-Time System

A hard real-time system is one in which failure to meet even a single deadline may lead to complete or catastrophic system failure.

Soft Real-Time System

A firm real-time system is one in which a few missed deadlines will not lead to total failure, but missing more than a few may lead to complete or catastrophic system failure

Degrees of “Real-Time”

Soft Real-Time System

A soft real-time system is one in which performance is degraded but not destroyed by failure to meet response-time constraints

Hard Real-Time System

A hard real-time system is one in which failure to meet even a single deadline may lead to complete or catastrophic system failure.

Soft Real-Time System

A firm real-time system is one in which a few missed deadlines will not lead to total failure, but missing more than a few may lead to complete or catastrophic system failure

Degrees of “Real-Time”—Examples

- Hard real-time software systems have a set of strict deadlines, and missing a deadline is considered a system failure —airplane sensor and autopilot systems, spacecrafts and planetary rovers.
- Firm real-time systems treat information delivered/computations made after a deadline as invalid. Like soft real-time systems, they do not fail after a missed deadline, and they may degrade QoS if a deadline is missed —financial forecast systems, robotic assembly lines.
- Soft real-time systems try to reach deadlines but do not fail if a deadline is missed. However, they may degrade their quality of service in such an event to improve responsiveness. —audio and video delivery software for entertainment (lag is undesirable but not catastrophic).

Degrees of “Real-Time”—Examples

- Hard real-time software systems have a set of strict deadlines, and missing a deadline is considered a system failure —airplane sensor and autopilot systems, spacecrafts and planetary rovers.
- Firm real-time systems treat information delivered/computations made after a deadline as invalid. Like soft real-time systems, they do not fail after a missed deadline, and they may degrade QoS if a deadline is missed —financial forecast systems, robotic assembly lines.
- Soft real-time systems try to reach deadlines but do not fail if a deadline is missed. However, they may degrade their quality of service in such an event to improve responsiveness. —audio and video delivery software for entertainment (lag is undesirable but not catastrophic).

Degrees of “Real-Time”—Examples

- Hard real-time software systems have a set of strict deadlines, and missing a deadline is considered a system failure —airplane sensor and autopilot systems, spacecrafts and planetary rovers.
- Firm real-time systems treat information delivered/computations made after a deadline as invalid. Like soft real-time systems, they do not fail after a missed deadline, and they may degrade QoS if a deadline is missed —financial forecast systems, robotic assembly lines.
- Soft real-time systems try to reach deadlines but do not fail if a deadline is missed. However, they may degrade their quality of service in such an event to improve responsiveness. —audio and video delivery software for entertainment (lag is undesirable but not catastrophic).

Real-Time Punctuality

Real-time punctuality means that every response time has an average value, t_R , with upper and lower bounds of $t_R + \epsilon_U$ and $t_R - \epsilon_L$, respectively, and $\epsilon_U, \epsilon_L \rightarrow 0^+$

- In all practical systems, the values of ϵ_U and ϵ_L are nonzero, though they may be very small
- The nonzero values are due to cumulative latency and propagation-delay components (hardware/software)
- Such response times contain jitter¹ within the interval $t \in [-\epsilon_L, +\epsilon_U]$

¹jitter is the amount of variation in response time

Where Do Deadlines Come from?

- Deadlines are based on the underlying physical phenomena of the system under control
- Punctuality is another measure related to response times for example in periodically sampled systems with high sampling rates (e.g., in audio and video signal processing)
- Examples:
 - An elevator door is automatically operated and it may have a capacitive safety edge for sensing possible passengers between the closing door blades
 - Thus, the door blades can be quickly reopened before they touch the passenger and cause discomfort or even threaten the passenger's safety
 - What is the required system response time from when it recognizes that a passenger is between the closing door blades and starting to reopen the door?

Where Do Deadlines Come from?

This response time consists of five independent latency components:

- 1 Sensor² ($t_{SE_{min}}=5ms$, $t_{SE_{max}}=15ms$)
- 2 Hardware ($t_{HW_{min}}=1\mu s$, $t_{HW_{max}}=2\mu s$)
- 3 System software³ ($t_{SS_{min}}=16\mu s$, $t_{SS_{max}}=48\mu s$)
- 4 Application software⁴ ($t_{AS_{min}}=0.5\mu s$, $t_{HW_{max}}=0.5\mu s$)
- 5 Door drive⁵ ($t_{DD_{min}}=300ms$, $t_{HW_{max}}=500ms$)

Consequently:

- minimum response time— $t_{min} = 305ms$
- maximum response time— $t_{max} = 515ms$

² Infrared beam sensors used to detect a person and prevent the doors from closing if a person or an object blocks the doorway

³ System Software is a set of programs that control and manage the operations of computer hardware. It also helps application programs to execute correctly. Example: Operating system, programming language, Communication software, etc.

⁴ Application Software acts as a mediator between the end-user and System Software. It is a program that does real work for the user. It is mostly created to perform a specific task for a user.

⁵ used for one side opening and central opening of the elevator's door

Where Do Deadlines Come from?

This response time consists of five independent latency components:

- 1 Sensor² ($t_{SE_{min}}=5ms$, $t_{SE_{max}}=15ms$)
- 2 Hardware ($t_{HW_{min}}=1\mu s$, $t_{HW_{max}}=2\mu s$)
- 3 System software³ ($t_{SS_{min}}=16\mu s$, $t_{SS_{max}}=48\mu s$)
- 4 Application software⁴ ($t_{AS_{min}}=0.5\mu s$, $t_{HW_{max}}=0.5\mu s$)
- 5 Door drive⁵ ($t_{DD_{min}}=300ms$, $t_{HW_{max}}=500ms$)

Consequently:

- minimum response time— $t_{min} = 305ms$
- maximum response time— $t_{max} = 515ms$

² Infrared beam sensors used to detect a person and prevent the doors from closing if a person or an object blocks the doorway

³ System Software is a set of programs that control and manage the operations of computer hardware. It also helps application programs to execute correctly. Example: Operating system, programming language, Communication software, etc.

⁴ Application Software acts as a mediator between the end-user and System Software. It is a program that does real work for the user. It is mostly created to perform a specific task for a user.

⁵ used for one side opening and central opening of the elevator's door

Where Do Deadlines Come from?

This response time consists of five independent latency components:

- 1 Sensor² ($t_{SE_{min}}=5ms$, $t_{SE_{max}}=15ms$)
- 2 Hardware ($t_{HW_{min}}=1\mu s$, $t_{HW_{max}}=2\mu s$)
- 3 System software³ ($t_{SS_{min}}=16\mu s$, $t_{SS_{max}}=48\mu s$)
- 4 Application software⁴ ($t_{AS_{min}}=0.5\mu s$, $t_{HW_{max}}=0.5\mu s$)
- 5 Door drive⁵ ($t_{DD_{min}}=300ms$, $t_{HW_{max}}=500ms$)

Consequently:

- minimum response time— $t_{min} = 305ms$
- maximum response time— $t_{max} = 515ms$

- ² Infrared beam sensors used to detect a person and prevent the doors from closing if a person or an object blocks the doorway
- ³ System Software is a set of programs that control and manage the operations of computer hardware. It also helps application programs to execute correctly. Example: Operating system, programming language, Communication software, etc.
- ⁴ Application Software acts as a mediator between the end-user and System Software. It is a program that does real work for the user. It is mostly created to perform a specific task for a user.
- ⁵ used for one side opening and central opening of the elevator's door

Where Do Deadlines Come from?

This response time consists of five independent latency components:

- 1 Sensor² ($t_{SE_{min}}=5ms$, $t_{SE_{max}}=15ms$)
- 2 Hardware ($t_{HW_{min}}=1\mu s$, $t_{HW_{max}}=2\mu s$)
- 3 System software³ ($t_{SS_{min}}=16\mu s$, $t_{SS_{max}}=48\mu s$)
- 4 Application software⁴ ($t_{AS_{min}}=0.5\mu s$, $t_{HW_{max}}=0.5\mu s$)
- 5 Door drive⁵ ($t_{DD_{min}}=300ms$, $t_{HW_{max}}=500ms$)

Consequently:

- minimum response time— $t_{min} = 305ms$
- maximum response time— $t_{max} = 515ms$

- ² Infrared beam sensors used to detect a person and prevent the doors from closing if a person or an object blocks the doorway
- ³ System Software is a set of programs that control and manage the operations of computer hardware. It also helps application programs to execute correctly. Example: Operating system, programming language, Communication software, etc.
- ⁴ Application Software acts as a mediator between the end-user and System Software. It is a program that does real work for the user. It is mostly created to perform a specific task for a user.
- ⁵ used for one side opening and central opening of the elevator's door

Where Do Deadlines Come from?

This response time consists of five independent latency components:

- 1 Sensor² ($t_{SE_{min}}=5ms$, $t_{SE_{max}}=15ms$)
- 2 Hardware ($t_{HW_{min}}=1\mu s$, $t_{HW_{max}}=2\mu s$)
- 3 System software³ ($t_{SS_{min}}=16\mu s$, $t_{SS_{max}}=48\mu s$)
- 4 Application software⁴ ($t_{AS_{min}}=0.5\mu s$, $t_{HW_{max}}=0.5\mu s$)
- 5 Door drive⁵ ($t_{DD_{min}}=300ms$, $t_{HW_{max}}=500ms$)

Consequently:

- minimum response time— $t_{min} = 305ms$
- maximum response time— $t_{max} = 515ms$

- ² Infrared beam sensors used to detect a person and prevent the doors from closing if a person or an object blocks the doorway
- ³ System Software is a set of programs that control and manage the operations of computer hardware. It also helps application programs to execute correctly. Example: Operating system, programming language, Communication software, etc.
- ⁴ Application Software acts as a mediator between the end-user and System Software. It is a program that does real work for the user. It is mostly created to perform a specific task for a user.
- ⁵ used for one side opening and central opening of the elevator's door

Where Do Deadlines Come from?

This response time consists of five independent latency components:

- 1 Sensor² ($t_{SE_{min}}=5ms$, $t_{SE_{max}}=15ms$)
- 2 Hardware ($t_{HW_{min}}=1\mu s$, $t_{HW_{max}}=2\mu s$)
- 3 System software³ ($t_{SS_{min}}=16\mu s$, $t_{SS_{max}}=48\mu s$)
- 4 Application software⁴ ($t_{AS_{min}}=0.5\mu s$, $t_{HW_{max}}=0.5\mu s$)
- 5 Door drive⁵ ($t_{DD_{min}}=300ms$, $t_{HW_{max}}=500ms$)

Consequently:

- minimum response time— $t_{min} = 305ms$
- maximum response time— $t_{max} = 515ms$

- ² Infrared beam sensors used to detect a person and prevent the doors from closing if a person or an object blocks the doorway
- ³ System Software is a set of programs that control and manage the operations of computer hardware. It also helps application programs to execute correctly. Example: Operating system, programming language, Communication software, etc.
- ⁴ Application Software acts as a mediator between the end-user and System Software. It is a program that does real work for the user. It is mostly created to perform a specific task for a user.
- ⁵ used for one side opening and central opening of the elevator's door

Where Do Deadlines Come from?

This response time consists of five independent latency components:

- 1 Sensor² ($t_{SE_{min}}=5ms$, $t_{SE_{max}}=15ms$)
- 2 Hardware ($t_{HW_{min}}=1\mu s$, $t_{HW_{max}}=2\mu s$)
- 3 System software³ ($t_{SS_{min}}=16\mu s$, $t_{SS_{max}}=48\mu s$)
- 4 Application software⁴ ($t_{AS_{min}}=0.5\mu s$, $t_{HW_{max}}=0.5\mu s$)
- 5 Door drive⁵ ($t_{DD_{min}}=300ms$, $t_{HW_{max}}=500ms$)

Consequently:

■ minimum response time— $t_{min} = 305ms$

■ maximum response time— $t_{max} = 515ms$

² Infrared beam sensors used to detect a person and prevent the doors from closing if a person or an object blocks the doorway

³ System Software is a set of programs that control and manage the operations of computer hardware. It also helps application programs to execute correctly. Example: Operating system, programming language, Communication software, etc.

⁴ Application Software acts as a mediator between the end-user and System Software. It is a program that does real work for the user. It is mostly created to perform a specific task for a user.

⁵ used for one side opening and central opening of the elevator's door

Where Do Deadlines Come from?

This response time consists of five independent latency components:

- 1 Sensor² ($t_{SE_{min}}=5ms$, $t_{SE_{max}}=15ms$)
- 2 Hardware ($t_{HW_{min}}=1\mu s$, $t_{HW_{max}}=2\mu s$)
- 3 System software³ ($t_{SS_{min}}=16\mu s$, $t_{SS_{max}}=48\mu s$)
- 4 Application software⁴ ($t_{AS_{min}}=0.5\mu s$, $t_{HW_{max}}=0.5\mu s$)
- 5 Door drive⁵ ($t_{DD_{min}}=300ms$, $t_{HW_{max}}=500ms$)

Consequently:

- minimum response time— $t_{min} = 305ms$
- maximum response time— $t_{max} = 515ms$

² Infrared beam sensors used to detect a person and prevent the doors from closing if a person or an object blocks the doorway

³ System Software is a set of programs that control and manage the operations of computer hardware. It also helps application programs to execute correctly. Example: Operating system, programming language, Communication software, etc.

⁴ Application Software acts as a mediator between the end-user and System Software. It is a program that does real work for the user. It is mostly created to perform a specific task for a user.

⁵ used for one side opening and central opening of the elevator's door

Event

Any occurrence that causes the program counter to change non-sequentially is considered a change of flow-of-control, and thus an event

Release Time

The release time is the time at which an instance of a scheduled task is ready to run, and is generally associated with an interrupt

Taxonomy⁶ of Events

- An event can be either synchronous or asynchronous
 - **Synchronous** events are those that occur at predictable times in the flow-of-control
 - **Asynchronous** events occur at unpredictable points in the flow-of-control and are usually caused by external sources
- Moreover, events can be periodic, aperiodic or sporadic
 - A real-time clock that pulses regularly is a **periodic event**
 - Events that do not occur at regular periods are called **aperiodic**
 - Aperiodic events that tend to occur very infrequently are called **sporadic**

TAB 1. Various Types of Events

Type	Periodic	Aperiodic	Sporadic
Synchronous	Cyclic code	Conditional branch	Divide-by-zero (trap) interrupt
Asynchronous	Clock interrupt	Regular, but not fixed-period interrupt	Power-loss alarm

⁶Taxonomy is the practice and science of categorization or classification

Deterministic systems

- For any physical system, certain states exist under which the system is considered to be out of control
- The software controlling such a system must therefore avoid these states
- In embedded real-time systems, maintaining overall control is extremely important
- Software control of any real-time system and associated hardware is maintained when the next state of the system, given the current state and a set of inputs, is predictable

Deterministic system

A system is deterministic, if for each possible state and each set of inputs, a unique set of outputs and next state of the system can be determined

Deterministic systems

- For any physical system, certain states exist under which the system is considered to be out of control
- The software controlling such a system must therefore avoid these states
- In embedded real-time systems, maintaining overall control is extremely important
- Software control of any real-time system and associated hardware is maintained when the next state of the system, given the current state and a set of inputs, is predictable

Deterministic system

A system is deterministic, if for each possible state and each set of inputs, a unique set of outputs and next state of the system can be determined

Deterministic systems

- Event determinism means the next states and outputs of a system are known for each set of inputs that trigger events
- Thus, a system that is deterministic is also event deterministic
- However, event determinism may not imply determinism
- While it is a significant challenge to design systems that are completely event deterministic, it is possible to inadvertently end up with a system that is non-deterministic
- Finally, if in a deterministic system the response time for each set of outputs is known, then, the system also exhibits temporal determinism
- A side benefit of designing deterministic systems is that guarantees can be given that the system will be able to respond at any time, and in the case of temporally deterministic systems, when they will respond

CPU Utilization

- The final term to be defined is a critical measure of real-time system performance
- Because the CPU continues to execute instructions as long as power is applied, it will more or less frequently execute instructions that are not related to the fulfillment of a specific deadline
- The measure of the relative time spent doing non-idle processing indicates how much real-time processing is occurring

CPU Utilization Factor

The CPU utilization or time-loading factor, U , is a relative measure of the non-idle processing taking place

CPU Utilization

- The final term to be defined is a critical measure of real-time system performance
- Because the CPU continues to execute instructions as long as power is applied, it will more or less frequently execute instructions that are not related to the fulfillment of a specific deadline
- The measure of the relative time spent doing non-idle processing indicates how much real-time processing is occurring

CPU Utilization Factor

The CPU utilization or time-loading factor, U , is a relative measure of the non-idle processing taking place

CPU Utilization Zones

- A system is said to be time-overloaded if $U > 100\%$
- Systems that are too highly utilized are problematic —additions, changes, or corrections cannot be made to the system without risk of time-overloading
- On the other hand, systems that are not sufficiently utilized are not necessarily cost-effective —The system was over-engineered and that costs could likely be reduced with less expensive hardware
- While a utilization of 50% is common for new products, 80% might be acceptable for systems that do not expect growth
- However, 70% as a target for U is one of the potentially useful results in the theory of real-time systems where tasks are periodic and independent

TAB 2. CPU Utilization Zones

CPU utilization [in %]	Zone Type	Typical Application
< 26	Unnecessarily safe	Various
26 – 50	Very safe	Various
51 – 68	Safe	Various
69	Theoretical limit	Embedded systems
70 – 82	Questionable	Embedded systems
83 – 99	Dangerous	Embedded systems
100	Critical	Marginally stressed system
> 100	Overloaded	Stressed system

Calculation of the CPU utilization

- U is calculated by summing the contribution of utilization factors for each task
- Suppose a system has $n \geq 1$ periodic tasks, each with an execution period of p_i
- If task i is known to have a worst-case execution time of e_i , then the utilization factor, u_i , for task i is given in Equation (1)

$$u_i = \frac{e_i}{p_i} \quad (1)$$

- The overall system utilization factor, U , is given by Equation (2)

$$\begin{aligned} U &= \sum_{i=1}^n u_i \\ &= \sum_{i=1}^n \frac{e_i}{p_i} \end{aligned} \quad (2)$$

- In practice, the determination of e_i can be difficult, in which case estimation or measuring must be used
- For aperiodic and sporadic tasks, u_i is calculated by assuming a worst-case execution period

Example—calculation of CPU utilization

An individual elevator controller in a bank of elevators has the following tasks with execution periods of p_i and worst-case execution times of $e_i \in [1, 2, 3, 4]$ as shown in **Table 3**:

- Task 1—Communicate with the group dispatcher.
- Task 2—Update the car position information and manage floor-to-floor runs as well as door control.
- Task 3—Register and cancel car calls.
- Task 4—Miscellaneous system supervisions.

The CPU utilization is computed as **Equation (3)**

$$U = \sum_{i=1}^n \frac{e_i}{p_i} \quad (3)$$

We note that $U = 31\%$ —which is a very safe zone

TAB 3. CPU utilization

i	e_i	p_i	$\frac{e_i}{p_i}$
1	17	500	0.03
2	4	25	0.16
3	1	75	0.01
4	20	200	0.10
U			0.31%

Usual Misconceptions

- **Real-time systems are synonymous with “fast” systems** —Many (but not all) hard real-time systems deal with deadlines in the tens of milliseconds
- **There are universal, widely accepted methodologies for real-time systems specification and design** —There is still no methodology available that answers all of the challenges of real-time specification and design all the time and for all applications
- **There is no more a need to build a real-time operating system, because many commercial products exist**—Commercial solutions have certainly their place, but choosing when to use an off-the-shelf solution and choosing the right one are continuing challenges
- **Rate-monotonic analysis has solved “the real-time problem”** —While rate-monotonic systems provide guidance in the design of real-time systems, and while there is theory surrounding them, they are not a panacea
- **The study of real-time systems is mostly about scheduling theory**—While it is scholarly to study scheduling theory, most published results require impractical simplifications and clairvoyance in order to make the theory work

Usual Misconceptions

- **Real-time systems are synonymous with “fast” systems** —Many (but not all) hard real-time systems deal with deadlines in the tens of milliseconds
- **There are universal, widely accepted methodologies for real-time systems specification and design** —There is still no methodology available that answers all of the challenges of real-time specification and design all the time and for all applications
- **There is no more a need to build a real-time operating system, because many commercial products exist**—Commercial solutions have certainly their place, but choosing when to use an off-the-shelf solution and choosing the right one are continuing challenges
- **Rate-monotonic analysis has solved “the real-time problem”** —While rate-monotonic systems provide guidance in the design of real-time systems, and while there is theory surrounding them, they are not a panacea
- **The study of real-time systems is mostly about scheduling theory**—While it is scholarly to study scheduling theory, most published results require impractical simplifications and clairvoyance in order to make the theory work

Usual Misconceptions

- **Real-time systems are synonymous with “fast” systems** —Many (but not all) hard real-time systems deal with deadlines in the tens of milliseconds
- **There are universal, widely accepted methodologies for real-time systems specification and design** —There is still no methodology available that answers all of the challenges of real-time specification and design all the time and for all applications
- **There is no more a need to build a real-time operating system, because many commercial products exist**—Commercial solutions have certainly their place, but choosing when to use an off-the-shelf solution and choosing the right one are continuing challenges
- **Rate-monotonic analysis has solved “the real-time problem”** —While rate-monotonic systems provide guidance in the design of real-time systems, and while there is theory surrounding them, they are not a panacea
- **The study of real-time systems is mostly about scheduling theory**—While it is scholarly to study scheduling theory, most published results require impractical simplifications and clairvoyance in order to make the theory work

Usual Misconceptions

- **Real-time systems are synonymous with “fast” systems** —Many (but not all) hard real-time systems deal with deadlines in the tens of milliseconds
- **There are universal, widely accepted methodologies for real-time systems specification and design** —There is still no methodology available that answers all of the challenges of real-time specification and design all the time and for all applications
- **There is no more a need to build a real-time operating system, because many commercial products exist**—Commercial solutions have certainly their place, but choosing when to use an off-the-shelf solution and choosing the right one are continuing challenges
- **Rate-monotonic analysis has solved “the real-time problem”** —While rate-monotonic systems provide guidance in the design of real-time systems, and while there is theory surrounding them, they are not a panacea
- **The study of real-time systems is mostly about scheduling theory**—While it is scholarly to study scheduling theory, most published results require impractical simplifications and clairvoyance in order to make the theory work

Usual Misconceptions

- **Real-time systems are synonymous with “fast” systems** —Many (but not all) hard real-time systems deal with deadlines in the tens of milliseconds
- **There are universal, widely accepted methodologies for real-time systems specification and design** —There is still no methodology available that answers all of the challenges of real-time specification and design all the time and for all applications
- **There is no more a need to build a real-time operating system, because many commercial products exist**—Commercial solutions have certainly their place, but choosing when to use an off-the-shelf solution and choosing the right one are continuing challenges
- **Rate-monotonic analysis has solved “the real-time problem”** —While rate-monotonic systems provide guidance in the design of real-time systems, and while there is theory surrounding them, they are not a panacea
- **The study of real-time systems is mostly about scheduling theory**—While it is scholarly to study scheduling theory, most published results require impractical simplifications and clairvoyance in order to make the theory work

Multidisciplinary Design Challenges

- Real-time systems is a truly multi-dimensional subdiscipline of computer systems engineering
- Therefore, it stands out as a fascinating study area with a versatile set of design challenges
- The fundamentals of real-time systems are well established and have considerable permanence
- Still, real-time systems is a lively developing area

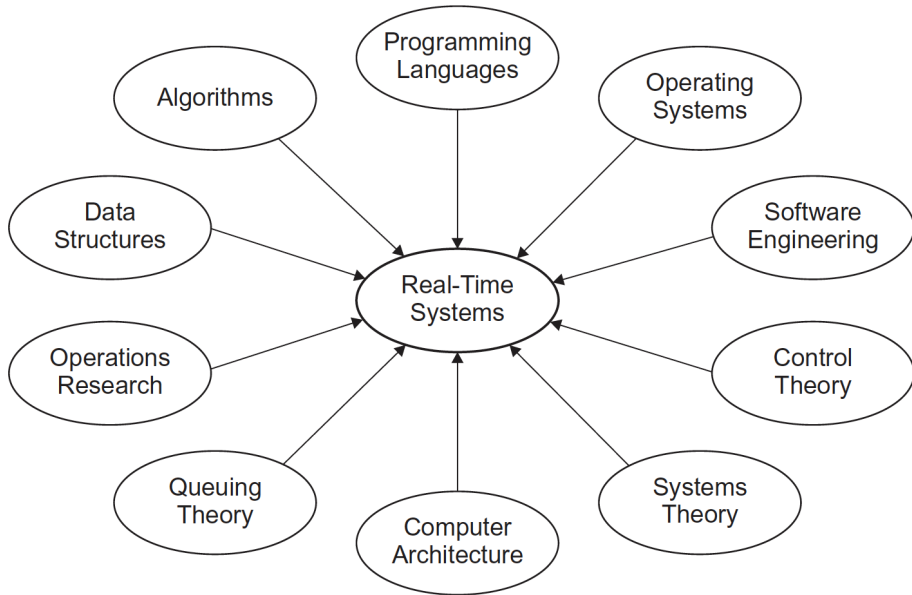


FIG 4. A variety of disciplines that affect real-time systems engineering

Practical Real-Time Embedded Systems

■ Aerospace

- Flight control
- Navigation
- Pilot interface

■ Automotive

- Airbag deployment
- Antilock braking
- Fuel injection

■ Household

- Microwave oven
- Rice cooker
- Washing machine

■ Industrial

- Crane
- Paper machine
- Welding robot

■ Multimedia

- Console game
- Home theater
- Simulator

■ Medical

- Intensive care monitor
- Magnetic resonance imaging
- Remote surgery

Impact on System Architecture

- Must avoid non-determinism in the system.
- Direct Memory Access (DMA) by peripheral devices—Contention for system bus.
- Cache
- Interrupts generated by I/O devices.
- Memory Management (paging)
- Dynamic data structures, recursion, unbounded loop —on the programming language level
- Nondeterministic algorithm⁷
- Bugs in the software e.g., memory leaks, using uninitialized memory, critical sections access violation, unclosed resources, race conditions,

⁷https://en.wikipedia.org/wiki/Nondeterministic_algorithm