

# Scheduling of Dependent Tasks

Kizito NKURIKIYEYU,

Ph.D.

## Readings

- Read Chapter 3 of Cottet et al. (2002). Scheduling in Real-Time Systems.

### Topics

- Task precedence relationships
- Sharing critical resources
- Mutual exclusion
- Priority inversion
- Deadlock

**SCHEDULING  
IN REAL-TIME  
SYSTEMS**

Frédéric Cottet | Judith Delacroix | Claude Kuhn | Zaid Mameri



<sup>1</sup>Readings are based on Cottet, F., Delacroix, J., Mameri, Z., & Kaiser, C. (2002). Scheduling in Real-Time Systems. Wiley.

Kizito NKURIKIYEYU, Ph.D.

Scheduling of Dependent Tasks

October 15, 2021


1 / 25

## Introduction

- The previous lecture assumed tasks were **independent**, i.e., there was no relationship between them
- This is too **simplistic** and does not reflect reality
- In most real-world application, inter-task cooperation and inter-task dependencies are a must
  - some tasks must respect the processing order
  - mutual exclusion to protect shared resources
  - precedence constraints that correspond to synchronization or communication among tasks

## Tasks with Precedence Relationships

- **precedence constraint** between two tasks  $\tau_i$  and  $\tau_j$  is denoted as  $\tau_i \rightarrow \tau_j$  if the execution of task  $\tau_i$  precedes that of task  $\tau_j$ .
- In this case, task  $\tau_j$  must await the completion of task  $\tau_i$  before it can execute



**FIG 1. Example of two precedence graphs related to a set of nine tasks**  
The relationships is described through a graph where the nodes represent tasks and the arrows express the precedence constraint between two nodes.

# Tasks with Precedence Relationships

- The previous precedence acyclic graph, however, represents a partial order on the task set.
- In general, we consider cases where  $n$  successive instance of a task can precede one instance of another task or vice versa.
- Fig. 2 shows an example of a generalized precedence relationship where the rate of communicating task are not equal.




FIG 2. Example of a generalized precedence relationship between two task

# Tasks with Precedence Relationships

if  $\tau_i \rightarrow \tau_j$ , then the task parameters must be in accordance with the following rules<sup>2</sup>:

- release times:  $r_j \geq r_i$
- priorities:  $priority_i \geq priority_j$ , in accordance with the scheduling algorithm

<sup>2</sup>Blazewicz, J. (1979). Deadline scheduling of tasks with ready times and resource constraints. Information Processing Letters, 8(2), 60–63. [https://doi.org/10.1016/0020-0190\(79\)90143-1](https://doi.org/10.1016/0020-0190(79)90143-1)

# Tasks with Precedence Relationships

Let's consider an example of in which  $\tau_i$  has to communicate its results to task  $\tau_j$

- $\tau_i$  and  $\tau_j$  have to be scheduled in a way that the execution of the  $k^{th}$  instance of task  $\tau_i$  precedes the the execution of the  $k^{th}$  instance of the task  $\tau_j$ . Thus, these task **have the same rate**, i.e.,  $T_i = T_j$
- $T_i \neq T_j$ , then tasks will run at the lowest rate sooner or later; consequently, the task with the shortest period will miss its deadline<sup>1</sup>.





FIG 3. Example of two precedence graphs related to a set of nine tasks.

# Precedence constraints and fixed-priority with rate monotonic algorithm

- We consider the **rate monotonic (RM)** and **deadline monotonic (DM)** algorithms
- In RM, tasks with shorter period get higher priorities.
- We want to respect this rule and figure out how to modify the task parameters in order to take account of precedence constraints, i.e. to obtain an independent task set with **modified parameters** with the following rules:
  - A task cannot start before its predecessors
  - A task cannot preempt its successors.
- If  $\tau_i \rightarrow \tau_j$ , then the release time and the priority of task parameters must be modified as follows:
  - $r_j^* \geq \max(r_j, r_i^*)$ , where  $r_i^*$  is the modified release time of task  $\tau_i$

## Example



**TAB 1.** Example of priority mapping taking care of precedence constraints and using the RM scheduling algorithm

Task	$\tau_1$	$\tau_2$	$\tau_3$	$\tau_4$	$\tau_5$	$\tau_6$
Priority	6	5	4	3	2	1

**FIG 4.** Precedence graphs of a set of six tasks

## Precedence constraints and fixed-priority with deadline monotonic algorithm

- With the deadline monotonic scheduling algorithm, tasks with shorter relative deadline get higher priorities
- The modifications of task parameters are close to those applied for RM scheduling except that the **relative deadline** is also changed in order to respect the priority assignment.
- If  $\tau_i \rightarrow \tau_j$ , then the release time, the relative deadline and the priority of the task parameters must be modified as follows:
  - $r_j^* \geq \max(r_j, r_i^*)$ , when  $r_i^*$  is the modified release time of task  $\tau_i$
  - $D_j^* \geq \max(D_j, D_i^*)$ , when  $D_i^*$  is the modified relative deadline of task  $\tau_i$
  - **priority**;  $\geq$  **priority**, in accordance with the DM scheduling algorithm

## Precedence constraints and the EDF algorithm

**review**—the **earliest deadline first (EDF)** algorithm assigns priority to tasks according to their absolute deadline: the task with the earliest deadline will be executed as the highest priority.


- with the EDF algorithm, the modification of task parameters relies on the deadline  $d$ .
- Rules for modifying release times and deadlines of tasks are based on the following observations<sup>3, 4</sup>:
  - 1 To get  $\tau_i \rightarrow \tau_j$ , the release time  $r_j^*$  of task  $\tau_j$  must be greater than or equal to its initial value or to the new release times  $\tau_i^*$  of its immediate predecessors  $\tau_i$  increased by their execution times  $C_i$

$$r_j^* \geq \max((r_i^* + C_i), r_j) \quad (1)$$

## Constraints and the EDF algorithm

- 2 If we have to get  $\tau_i \rightarrow \tau_j$ , the deadline  $d_j^*$  of task  $\tau_j$  has to be replaced by the minimum between its initial value  $d_j$  by the new deadline  $d_j^*$  of the immediate successors  $\tau_j$  decreased by their execution times  $C_j$  :

$$d_j^* \geq \min((d_j^* - C_j), d_i) \quad (2)$$



**FIG 5.** Modifications of task parameters in the case of EDF scheduling. The modifications begin with the tasks that have no predecessors for modifying their release times and with those with no successors for changing their deadlines. **Please see example on page 54.**

<sup>3</sup>Blazewicz J. (1997), Scheduling dependent tasks with different arrival times to meet deadlines, in Beilner H. and Gelenbe E. (eds) Modeling and Performance

# Tasks Sharing Critical Resources

## Resource Sharing

- example of **shared resource**—data structures (e.g., queue), variables, main memory area, file, set of registers, I/O unit, etc.
- Many shared resources do not allow simultaneous accesses but require **mutual exclusion**. These resources are called **exclusive resources**.
- No two tasks are allowed to operate on the resource at the same time.
- Protection methods: interrupt disabling<sup>5</sup> and using semaphore or mutex
- In FreeRTOS, The **taskENTER\_CRITICAL()** and **taskEXIT\_CRITICAL()** provide a basic critical section implementation that works by simply disabling interrupts, either globally, or up to a specific interrupt priority level.

```
1
2 taskENTER_CRITICAL();
```

Kizito NKURIKIYEYU, Ph.D.

Scheduling of Dependent Tasks

October 15, 2021

12 / 25

## Resource Sharing

- Task  $J_2$  has higher priority than task  $J_1$
- Task  $J_1$  is activated first and use the resource R (i.e, enters the critical section)
- If task  $J_2$  (with higher priority) tries access the processor, it will preempt task  $J_1$ . However, if it tries to access the shared resources, it is blocked due to the mutual exclusion guaranteed by the semaphore.
- When blocked, the task  $J_1$  can resume its execution and complete




FIG 6. Two tasks sharing one resource

## Mutual exclusion

In FreeRTOS, a mutex is a special type of semaphore that is used to **control access** to a resource that is shared between two or more tasks.

- When used in a mutual exclusion scenario, the mutex can be thought of as a **token** that is associated with the resource being shared.
- For a task to access the resource legitimately, it must first successfully **take** the token. When the token holder has finished with the resource, it must **give** the token back.
- Only when the token has been **returned** can another task successfully take the token, and then safely access the same shared resource.

```
1 SemaphoreHandle_t xMutex
2 int main( void) {
3     xMutex = xSemaphoreCreateMutex()
4     if (xMutex != NULL) {
```

```

1 void vTask1( void *pvParameters ) {
2     while(true) {
3         ...
4         xSemaphoreTake (xMutex,portMAX_DELAY);
5         /* access to exclusive resource */
6         xSemaphoreGive (xMutex)
7         ...
8     }
9 }
10 void vTask2( void *pvParameters ) {
11     while(true) {
12         ...
13         xSemaphoreTake (xMutex,portMAX_DELAY);
14         /* access to exclusive resource */
15         xSemaphoreGive (xMutex)
16         ...
17     }
18 }

```

## Priority inversion

- **Priority inversion** may occur in preemptive scheduling that is driven by fixed priority and where critical resources are protected by a mutual exclusion mechanism.
- **Priority inversion** —a case where a medium priority task is executed prior to a high priority task; this occurs because the latter is blocked—for an unbounded amount of time—by a low priority task. It is a consequence of shared resource access.
- Priority inversion, contravenes the scheduling specification and can induce deadline missing

## Priority inversion

Consider a task set composed of four tasks  $\tau_1, \tau_2, \tau_3, \tau_4$  having decreasing priorities (i.e.,  $\tau_1$  has the highest priority and  $\tau_4$  the lowest) and where Tasks  $\tau_2$  and  $\tau_4$  share a critical resource  $R_1$ , the access of which is **mutually exclusive**




FIG 7. Example of priority inversion phenomenon

## Priority inversion

- The lowest priority task  $\tau_4$  starts its execution first and after some time it enters a critical section using resource  $R_1$ .
- When task  $\tau_4$  is in its critical section, the higher priority task  $\tau_2$  is released and preempts task  $\tau_4$
- During the execution of task  $\tau_2$ , task  $\tau_3$  is released.
- Nevertheless, task  $\tau_3$ , having a lower priority than task  $\tau_2$ , must wait
- When task  $\tau_2$  needs to enter its **critical section**, associated with the critical resource  $R_1$  shared with task  $\tau_4$ , it finds that the corresponding resource  $R_1$  is held by task  $\tau_4$ .—Thus it is blocked
- The highest priority task able to execute is task  $\tau_3$ , So task  $\tau_3$ , gets the processor and runs.
- During this execution, the highest priority task  $\tau_1$  awakes. As a consequence task  $\tau_3$  is suspended and the processor is allocated to task  $\tau_1$ .

- At the end of execution of task  $\tau_1$ , task  $\tau_3$  can resume its execution until it reaches the end of its code.
- Now, only the **lowest priority task  $\tau_4$** , preempted in its critical section, can execute again. It resumes its execution until it releases critical resource  $R_1$  required by the higher priority task  $\tau_2$
- Then, task  $\tau_2$  can resume its execution by holding critical resource  $R_1$  necessary for its activity


■ **Remarks:**

- Task  $\tau_2$ 's maximum blocking time varies and depends on the duration of the critical section of the lower priority tasks sharing the resource with it (e.g.,  $\tau_2$ )
- The blocking time also depends on the execution time of the higher priority task  $\tau_1$
- A lower priority task,  $\tau_3$ , increased the blocking time of a higher priority task  $\tau_2$ , even if  $\tau_3$  does not share any critical resource with  $\tau_2$
- When there is priority inversion, the blocking time of each task

## Why this course?



**FIG 8.** Artist's conception of NASA's Mars Exploration Rover on Mars. Its mission almost failed due priority inversion.



**FIG 9.** Instrumentation of the Mars Rover

<sup>2</sup><http://www.cs.cornell.edu/courses/cs614/1999sp/papers/pathfinder.html>

## Mars rover and priority inversion

- A few days into the mission, the rover began experiencing total system resets, each resulting in losses of data<sup>2</sup>.
- Priority inversion was the root cause because VxWorks<sup>6</sup>'s preemptive priority scheduling
  - Its bus management task ran frequently with high priority and access to the bus was synchronized with mutual exclusion locks
  - The meteorological data gathering task ran a low priority thread and acquire a mutex when publishing its data, writes to the bus, and release the mutex
  - A communications task that ran with medium priority.
- It was possible for an interrupt to occur that caused the the **medium priority** communications task to be scheduled during the short interval while the **high priority** information bus thread was blocked waiting for the **low priority** meteorological data thread, consequently preventing the blocked information

## Solutions to Priority Inversion

- **Disallow preemption** during the execution of all critical sections.
  - simple approach
  - but it creates unnecessary blocking as unrelated tasks may be blocked.
- **Resource access protocols**—modify the priority of those tasks that cause blocking. When a task  $\tau_i$  blocks one or more higher priority tasks, it temporarily assumes a higher priority. Several approaches exist:
  - Priority Inheritance Protocol (PIP), for static priorities<sup>7, 8</sup>
  - Priority Ceiling Protocol (PCP), for static priorities<sup>9</sup>
  - Stack Resource Policy (SRP), for static and dynamic priorities<sup>10</sup>

<sup>7</sup><https://www.embedded.com/introduction-to-priority-inversion/>

<sup>8</sup><https://www.embedded.com/how-to-use-priority-inheritance/>

<sup>9</sup>[https://en.wikipedia.org/wiki/Priority\\_ceiling\\_protocol](https://en.wikipedia.org/wiki/Priority_ceiling_protocol)

<sup>10</sup>[https://en.wikipedia.org/wiki/Stack\\_Resource\\_Policy](https://en.wikipedia.org/wiki/Stack_Resource_Policy)

# Priority Inheritance Protocol

- **summary**—When a task  $\tau_i$  blocks one or more higher priority tasks, it temporarily assumes (inherits) the highest priority of the blocked tasks. It allows this task to use the critical resource as early as possible without going through the preemption. It avoids the unbounded priority inversion<sup>11</sup>.
- **assumptions**
  - $n$  tasks which cooperate through  $m$  shared resources
  - fixed priorities
  - all critical sections on a resource begin with a **take()** and end with a **give** operation
- **advantages**
  - It allows the different priority tasks to share the critical resources.
  - it avoids the unbounded priority inversion.
- **disadvantages**
  - can lead to **deadlock**<sup>12</sup>

# Deadlock phenomenon

**summary**—a situation in which two or more tasks are blocked indefinitely because each task is waiting for a resource acquired by another blocked task (Fig. 10).




FIG 10. Example of the deadlock phenomenon

- Two tasks  $\tau_1$  and  $\tau_2$  use two critical resources  $R_1$  and  $R_2$ .
- $\tau_1$  and  $\tau_2$  access  $R_1$  and  $R_2$  in reverse order. Moreover, the priority of task  $\tau_1$  is greater than that of task  $\tau_2$ .
- Now, suppose that task  $\tau_2$  executes first and locks resource  $R_1$ .

# Deadlock phenomenon

- During the critical section of task  $\tau_2$  using resource  $R_1$ , task  $\tau_1$  awakes and preempts task  $\tau_2$  before it can lock the second resource  $R_2$ .
- Task  $\tau_1$  needs resource  $R_2$  first, which is free, and it locks it.
- Then task  $\tau_1$  needs resource  $R_1$ , which is held by task  $\tau_2$ . So task  $\tau_2$  resumes and asks for resource  $R_2$ , which is not free.
- The final result is that task  $\tau_2$  is in possession of resource  $R_1$  but is waiting for resource  $R_2$  and task  $\tau_1$  is in possession of resource  $R_2$  but is waiting for resource  $R_1$ .
- Neither task  $\tau_1$  nor task  $\tau_2$  will release the resource until its pending request is satisfied.

The end